

AI SUMMER CAMP

TRANSFORMER & NLP

Content

1. Natural Language Processing over time
2. Transformer Motivation: How an RNN works
3. Transformer
 - a. Encoder-Decoder
 - b. Attention
 - c. Word Embedding
 - d. Positional Encoding
 - e. Dropout
 - f. Residual Connection
 - g. Advantages
4. Improved Transformers
5. Courses on NLP

Natural Language Processing

Common tasks include:

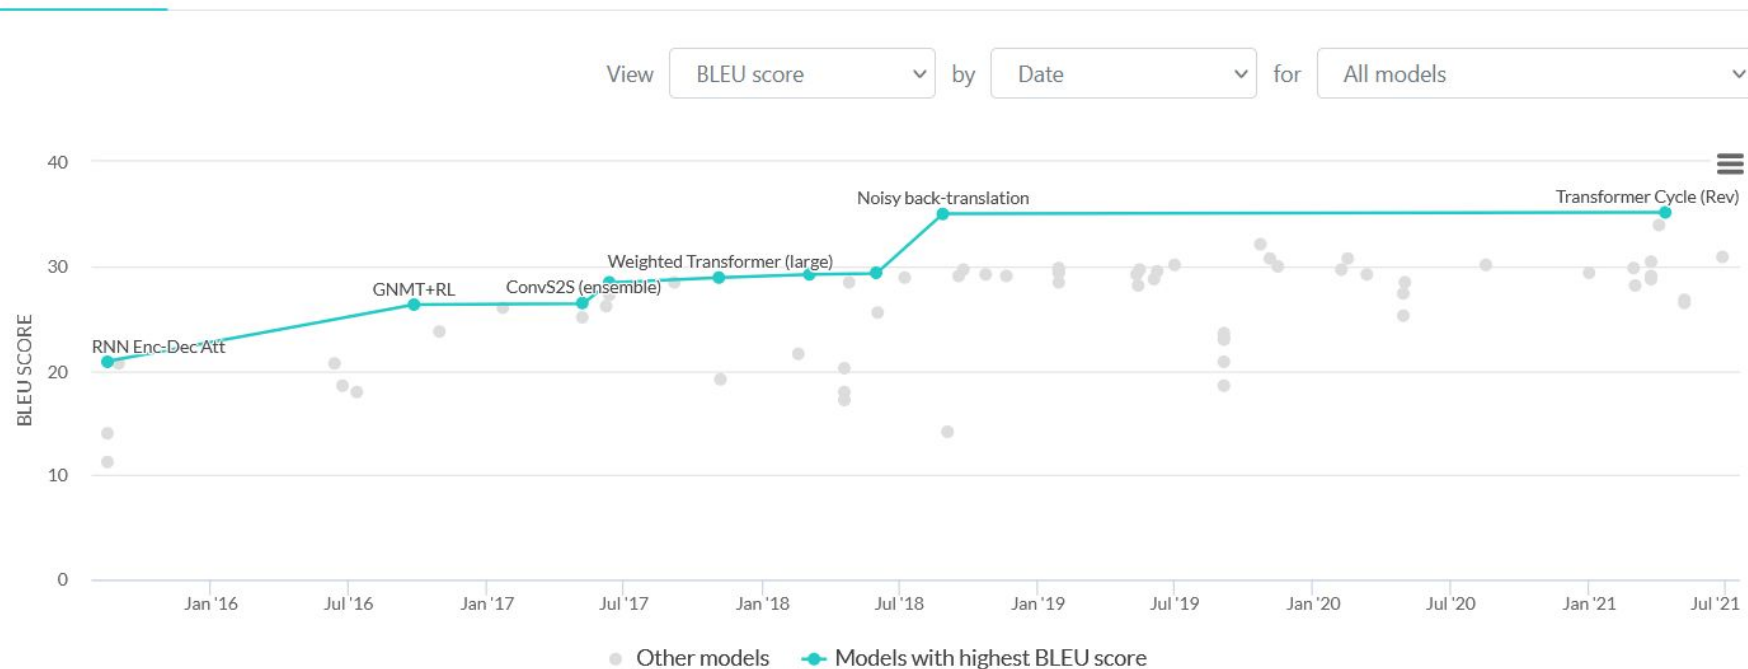
- text classification
- translation
- summarization
- named entity recognition
- dialogue (chatbots)
- question answering


















For more info visit <https://paperswithcode.com/area/natural-language-processing>

Machine Translation on WMT2014 English-German

Leaderboard

Dataset



1	Transformer Cycle (Rev)	35.14	33.54	✓	Lessons on Parameter Sharing across Layers in Transformers			2021	Transformer
2	Noisy back-translation	35.0	33.8	✓	Understanding Back-Translation at Scale			2018	
3	Transformer+Rep (Uni)	33.89	32.35	✓	Rethinking Perturbations in Encoder-Decoders for Fast Training			2021	Transformer
4	T5-11B	32.1		✓	Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer			2019	Transformer
5	Transformer + R-Drop	30.91		×	R-Drop: Regularized Dropout for Neural Networks			2021	Transformer
6	BERT-fused NMT	30.75		×	Incorporating BERT into Neural Machine Translation			2020	Transformer
7	Data Diversification - Transformer	30.7		×	Data Diversification: A Simple Strategy For Neural Machine Translation			2019	Transformer
8	Mask Attention Network (big)	30.4	215M	×	Mask Attention Networks: Rethinking and Strengthen Transformer			2021	
9	Transformer (ADMIN init)	30.1	29.5	×	Very Deep Transformers for Neural Machine Translation			2020	Transformer

Transformer

Attention Is All You Need

Ashish Vaswani* Google Brain avaswani@google.com	Noam Shazeer* Google Brain noam@google.com	Niki Parmar* Google Research nikip@google.com	Jakob Uszkoreit* Google Research usz@google.com
Llion Jones* Google Research llion@google.com	Aidan N. Gomez* † University of Toronto aidan@cs.toronto.edu	Lukasz Kaiser* Google Brain lukaszkaiser@google.com	
Illia Polosukhin* † illia.polosukhin@gmail.com			

Abstract

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks that include an encoder and a decoder. The best performing models also connect the encoder and decoder through an attention mechanism. We propose a new simple network architecture, the Transformer, based solely on attention mechanisms, dispensing with recurrence and convolutions entirely. Experiments on two machine translation tasks show these models to be superior in quality while being more parallelizable and requiring significantly less time to train. Our model achieves 28.4 BLEU on the WMT 2014 English-to-German translation task, improving over the existing best results, including ensembles, by over 2 BLEU. On the WMT 2014 English-to-French translation task, our model establishes a new single-model state-of-the-art BLEU score of 41.8 after training for 3.5 days on eight GPUs, a small fraction of the training costs of the best models from the literature. We show that the Transformer generalizes well to other tasks by applying it successfully to English constituency parsing both with large and limited training data.

Recurrent Neural Network (RNN)

The Vanilla RNN Model

First time-step ($t = 1$):

$$\mathbf{h}_1 = \tanh(W^{xh} \cdot \mathbf{x}_1 + W^{hh} \cdot \mathbf{h}_0)$$

$$\hat{\mathbf{y}}_1 = \text{softmax}(W^{hy} \cdot \mathbf{h}_1)$$

$$L_1 = CE(\hat{\mathbf{y}}_1, \mathbf{y}_1)$$

In general:


$$\mathbf{h}_t = \tanh(W^{xh} \cdot \mathbf{x}_t + W^{hh} \cdot \mathbf{h}_{t-1})$$

$$\hat{\mathbf{y}}_t = \text{softmax}(W^{hy} \cdot \mathbf{h}_t)$$


$$L_t = CE(\hat{\mathbf{y}}_t, \mathbf{y}_t)$$

In total:

$$L = \sum_t L_t$$

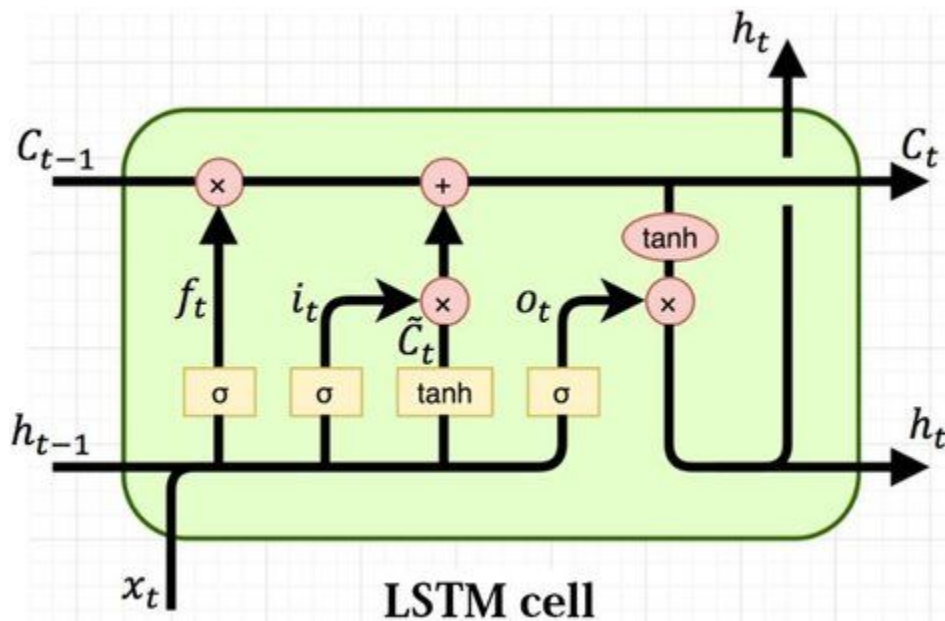


Backpropagation through time



$$\frac{\partial C_t}{\partial \mathbf{W}} = \sum_{t'=1}^t \frac{\partial C_t}{\partial h_t} \frac{\partial h_t}{\partial h_{t'}} \frac{\partial h_{t'}}{\partial \mathbf{W}}, \text{ where } \frac{\partial h_t}{\partial h_{t'}} = \prod_{k=t'+1}^t \frac{\partial h_k}{\partial h_{k-1}}$$

Long Short Term Memory (LSTM)



$$i_t = \sigma(x_t U^i + h_{t-1} W^i)$$

$$f_t = \sigma(x_t U^f + h_{t-1} W^f)$$

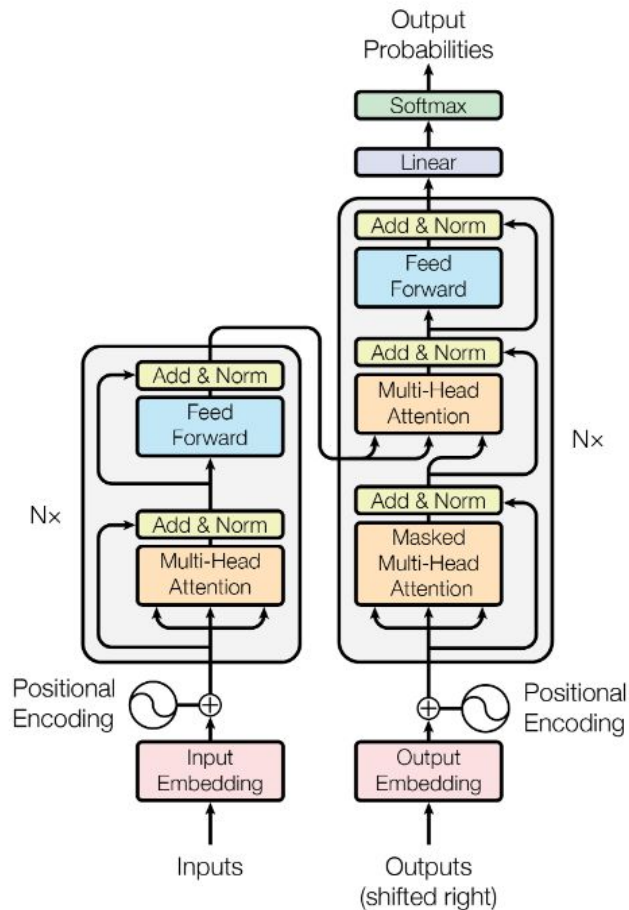
$$o_t = \sigma(x_t U^o + h_{t-1} W^o)$$

$$\tilde{C}_t = \tanh(x_t U^g + h_{t-1} W^g)$$

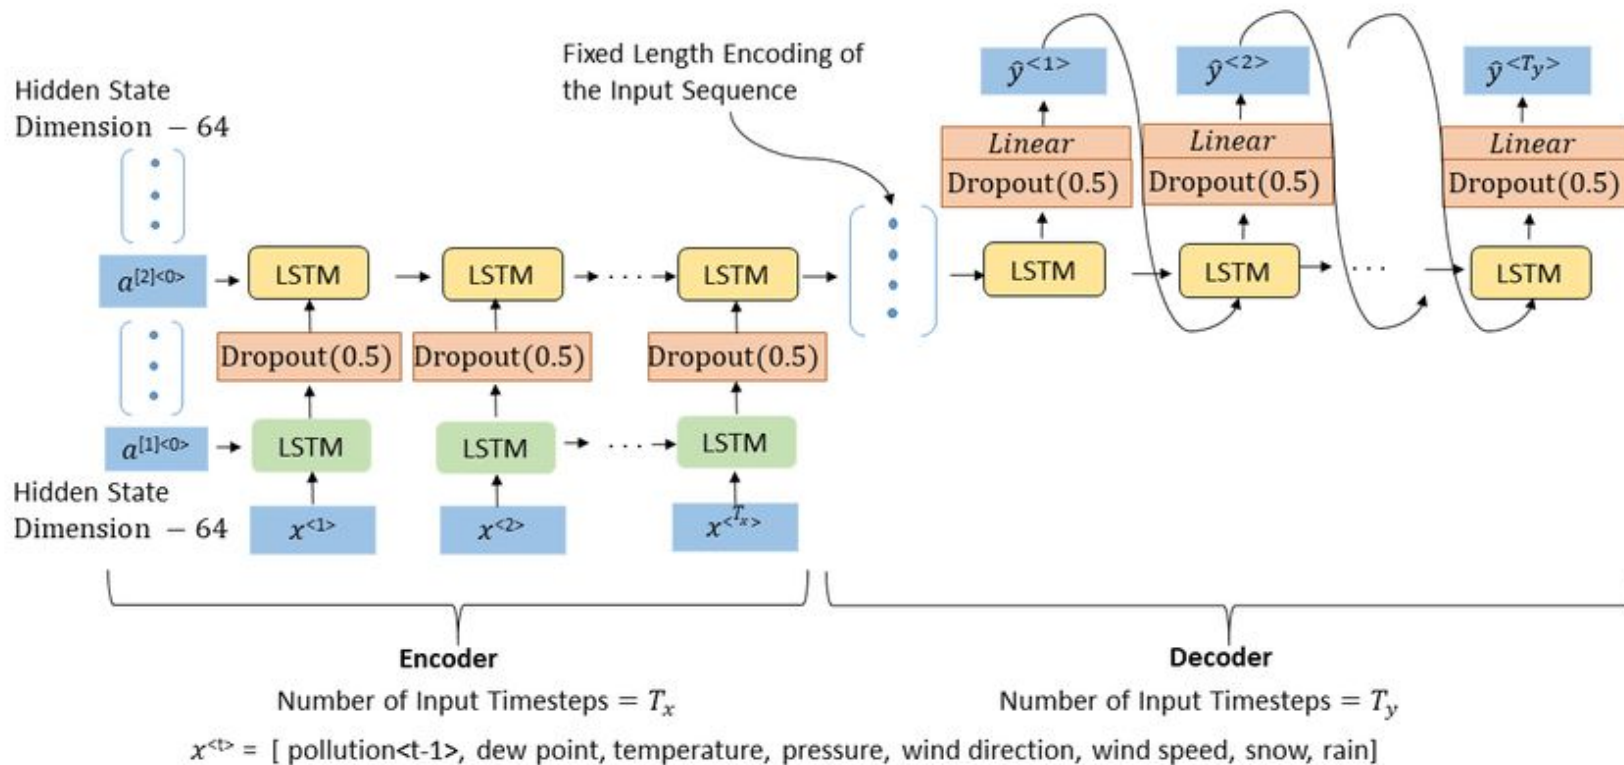
$$C_t = \sigma(f_t * C_{t-1} + i_t * \tilde{C}_t)$$

$$h_t = \tanh(C_t) * o_t$$


Transformer




Encoder Decoder




Word Embedding



Male-Female





Verb Tense



Country-Capital

<https://projector.tensorflow.org/>

Attention



Scaled Dot Product Attention

$$attention(Q,K,V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V \quad (4)$$

- Q -> query
- K -> key
- V -> value

yemek		koşarken		yemek		yemek					
<table border="1" style="background-color: #4a7ebb; color: white; width: 100px; height: 30px;">2</table> <table border="1" style="background-color: #4a7ebb; color: white; width: 100px; height: 30px;">-1</table>	●	<table border="1" style="background-color: #4a7ebb; color: white; width: 50px; height: 100px;"> <tr><td style="text-align: center;">-3</td></tr> <tr><td style="text-align: center;">1</td></tr> </table>	-3	1	= 2*(-3) + (-1)*1 = -7	<table border="1" style="background-color: #4a7ebb; color: white; width: 100px; height: 30px;">2</table> <table border="1" style="background-color: #4a7ebb; color: white; width: 100px; height: 30px;">-1</table>	●	<table border="1" style="background-color: #4a7ebb; color: white; width: 50px; height: 100px;"> <tr><td style="text-align: center;">2</td></tr> <tr><td style="text-align: center;">-1</td></tr> </table>	2	-1	= 5
-3											
1											
2											
-1											

(1)

yemek		yedik			
<table border="1" style="background-color: #4a7ebb; color: white; width: 100px; height: 30px;">2</table> <table border="1" style="background-color: #4a7ebb; color: white; width: 100px; height: 30px;">-1</table>	●	<table border="1" style="background-color: #4a7ebb; color: white; width: 50px; height: 100px;"> <tr><td style="text-align: center;">2</td></tr> <tr><td style="text-align: center;">1</td></tr> </table>	2	1	= 3
2					
1					

$$\textit{softmax}(-7) = 10^{-6}$$

$$\textit{softmax}(5) = 0.88 \quad (2)$$

$$\textit{softmax}(3) = 0.12$$

$$10^{-6} \begin{array}{|c|c|} \hline \text{koşarken} \\ \hline -3 & 1 \\ \hline \end{array} + 0.88 \begin{array}{|c|c|} \hline \text{yemek} \\ \hline 2 & -1 \\ \hline \end{array} + 0.12 \begin{array}{|c|c|} \hline \text{yedik} \\ \hline 2 & 1 \\ \hline \end{array}$$

(3)

$$= \begin{array}{|c|c|} \hline \text{yemek'} \\ \hline 2 & -0.76 \\ \hline \end{array}$$

$$\begin{array}{l} \text{koşarken} \\ \text{yemek} \\ \text{yedik} \end{array} \begin{array}{|c|c|} \hline \text{Query} \\ \hline \begin{array}{cc} -3 & 1 \\ 2 & -1 \\ 2 & 1 \end{array} \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline \text{Key}^T \\ \hline \begin{array}{ccc} -3 & 2 & 2 \\ 1 & -1 & 1 \end{array} \\ \hline \begin{array}{c} \text{koşarken} \\ \text{yemek} \\ \text{yedik} \end{array} \end{array} = \begin{array}{l} \text{koşarken} \\ \text{yemek} \\ \text{yedik} \end{array} \begin{array}{|c|c|c|} \hline \begin{array}{ccc} 10 & -7 & -5 \\ -7 & 5 & 3 \\ -5 & 3 & 5 \end{array} \\ \hline \begin{array}{c} \text{koşarken} \\ \text{yemek} \\ \text{yedik} \end{array} \end{array} \quad (5)$$

$$\text{softmax}\left(\frac{\begin{array}{|c|c|c|} \hline 10 & -7 & -5 \\ \hline -7 & 5 & 3 \\ \hline -5 & 3 & 5 \\ \hline \end{array}}{\sqrt{2}}\right) = \begin{array}{|c|c|c|} \hline 0.99 & 10^{-6} & 10^{-5} \\ \hline 10^{-4} & 0.80 & 0.19 \\ \hline 10^{-4} & 0.19 & 0.80 \\ \hline \end{array} \quad (6)$$

0.99	10^{-6}	10^{-5}
10^{-4}	0.80	0.19
10^{-4}	0.19	0.80

×

Value

-3	1
2	-1
2	1

=

Y

-2.97	0.99
1.98	-0.61
1.98	0.61

koşarken

yemek (7)


yedik

$$\begin{array}{c}
 \text{Word embedding} \\
 \begin{array}{c|c}
 \text{koşarken} & -3 & 1 \\
 \text{yemek} & 2 & -1 \\
 \text{yedik} & 2 & 1
 \end{array}
 \end{array}
 \begin{array}{c}
 \times \\
 \times \\
 \times
 \end{array}
 \begin{array}{c}
 \text{Weight}_Q \\
 \text{Weight}_K \\
 \text{Weight}_V
 \end{array}
 \begin{array}{c}
 = \\
 = \\
 =
 \end{array}
 \begin{array}{c}
 Q \\
 K \\
 V
 \end{array}
 \quad (8)$$

The diagram illustrates the matrix multiplication of word embeddings with weight matrices to produce query (Q), key (K), and value (V) matrices. The word embedding matrix is a 3x2 matrix with values:

koşarken	-3	1
yemek	2	-1
yedik	2	1

 Each of the three weight matrices (Weight_Q, Weight_K, Weight_V) is a 2x4 matrix of blue squares. The resulting Q, K, and V matrices are each 3x4 matrices of blue squares. The equation is labeled (8).



Positional Encoding

Solution #1

e_0	p_0		e_1	p_1		e_2	p_2		e_3	p_3
0.42	0		0.87	1		0.02	2		0.02	3
0.31	0		-0.64	1		0.01	2		0.01	3
0.73	0	+	0.81	1	+	-0.24	2	+	-0.24	3
0.36	0		0.26	1		-0.07	2		-0.07	3
0.99	0		-0.35	1		0.00	2		0.00	3

Positional Encoding

Solution #2


e_0	p_0		e_1	p_1		e_2	p_2		e_3	p_3
0.42	0		0.87	0.33		0.02	0.66		0.02	1
0.31	0		-0.64	0.33		0.01	0.66		0.01	1
0.73	0	+	0.81	0.33	+	-0.24	0.66	+	-0.24	1
0.36	0		0.26	0.33		-0.07	0.66		-0.07	1
0.99	0		-0.35	0.33		0.00	0.66		0.00	1

Positional Encoding

Solution #3

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$



<https://www.youtube.com/watch?v=dichIcUZfOw>

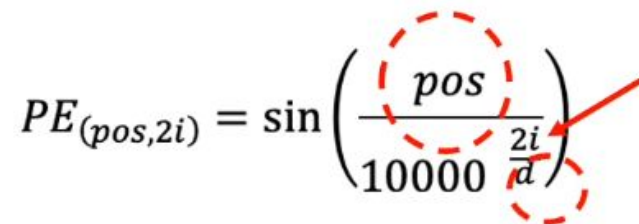


$$PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d}}}\right)$$



Diagram illustrating the positional encoding formula with annotations. The formula is $PE_{(pos, 2i)} = \sin\left(\frac{pos}{10000^{\frac{2i}{d}}}\right)$. Red dashed circles highlight the variables pos and $\frac{2i}{d}$. A red arrow points to the denominator $10000^{\frac{2i}{d}}$.



Dropout




(a) Standard Neural Net



(b) After applying dropout.

Residual Connection



Complexity

Table 1: Maximum path lengths, per-layer complexity and minimum number of sequential operations for different layer types. n is the sequence length, d is the representation dimension, k is the kernel size of convolutions and r the size of the neighborhood in restricted self-attention.

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$

Advantages

parallelizable (thus faster)

computationally less complex (most of the time)

better capture longer dependencies

more interpretable

Improved Transformers

- **BERT**
- GPT
- T5
- BART
- Pegasus
- XLM
- Reformer
- Longformer
- ELECTRA
- RoBERTa
- ...

Vision Transformers

AN IMAGE IS WORTH 16X16 WORDS: TRANSFORMERS FOR IMAGE RECOGNITION AT SCALE

<https://arxiv.org/pdf/2010.11929.pdf>



Courses & Resources

- Stanford University NLP w/ DL - <http://web.stanford.edu/class/cs224n/>
- Huggingface - <https://huggingface.co/course/chapter1>
- Deeplearning.ai NLP - <https://www.deeplearning.ai/program/natural-language-processing-specialization/>

Bibliography

- https://user.ceng.metu.edu.tr/~skalkan/DL/week_13.pdf
- https://kazemnejad.com/blog/transformer_architecture_positional_encoding/
- <https://www.youtube.com/watch?v=dichIcUZfOw>
- <https://arxiv.org/pdf/1706.03762.pdf>